



Chapter 6 – Cloud Resource Management and Scheduling

Resource management and scheduling

- Critical function of any man-made system.
- It affects the three basic criteria for the evaluation of a system:
 - Functionality.
 - Performance.
 - Cost.
- Scheduling in a computing system →→deciding how to allocate resources of a system, such as CPU cycles, memory, secondary storage space, I/O and network bandwidth, between users and tasks.
- Policies and mechanisms for resource allocation.
 - Policy →→principles guiding decisions.
 - Mechanisms →→the means to implement policies.

Motivation

- Cloud resource management .
 - ▣ Requires complex policies and decisions for multi-objective optimization.
 - ▣ It is challenging - the complexity of the system makes it impossible to have accurate global state information.
 - ▣ Affected by unpredictable interactions with the environment, e.g., system failures, attacks.
 - ▣ Cloud service providers are faced with large fluctuating loads which challenge the claim of cloud elasticity.
- The strategies for resource management for IaaS, PaaS, and SaaS are different.



Cloud resource management (CRM) policies

1. Admission control → prevent the system from accepting workload in violation of high-level system policies.
2. Capacity allocation → allocate resources for individual activations of a service.
3. Load balancing → distribute the workload evenly among the servers.
4. Energy optimization → minimization of energy consumption.
5. Quality of service (QoS) guarantees → ability to satisfy timing or other conditions specified by a Service Level Agreement.

Mechanisms for the implementation of resource management policies

- Control theory →→ uses the feedback to guarantee system stability and predict transient behavior.
- Machine learning →→ does not need a performance model of the system.
- Utility-based →→ require a performance model and a mechanism to correlate user-level performance with cost.
- Market-oriented/economic →→ do not require a model of the system, e.g., combinatorial auctions for bundles of resources.

Control theory application to cloud resource management (CRM)

- The main components of a control system:

- The inputs → the offered workload and the policies for admission control, the capacity allocation, the load balancing, the energy optimization, and the QoS guarantees in the cloud.
- The control system components → sensors used to estimate relevant measures of performance and *controllers* which implement various policies.
- The outputs → the resource allocations to the individual applications.

Feedback and Stability

- Control granularity → the level of detail of the information used to control the system.
 - Fine control → very detailed information about the parameters controlling the system state is used.
 - Coarse control → the accuracy of these parameters is traded for the efficiency of implementation.
- The controllers use the feedback provided by sensors to stabilize the system. Stability is related to the change of the output.
- Sources of instability in any control system:
 - The delay in getting the system reaction after a control action.
 - The granularity of the control, the fact that a small change enacted by the controllers leads to very large changes of the output.
 - Oscillations, when the changes of the input are too large and the control is too weak, such that the changes of the input propagate directly to the output.

Resource bundling

- ◆ Resources in a cloud are allocated in **bundles**.
- ◆ Users get maximum benefit from a specific combination of resources: CPU cycles, main memory, disk space, network bandwidth, and so on.
- ◆ Resource bundling complicates traditional resource allocation models and has generated an interest in economic models and, in particular, in **auction algorithms**.
- ◆ The bidding process aims to optimize an objective function $f(x,p)$.
- ◆ In the context of cloud computing, **an auction is the allocation of resources to the highest bidder**.

Cloud scheduling algorithms (1/2)

- Scheduling → responsible for resource sharing at several levels:
 - A server can be shared among several virtual machines.
 - A virtual machine could support several applications.
 - An application may consist of multiple threads.
- A scheduling algorithm should be **efficient**, **fair**, and **starvation-free**.
- The objectives of a scheduler:
 - Batch system → **maximize throughput** and **minimize turnaround time**.
 - Real-time system → **meet the deadlines** and **be predictable**.
- Best-effort: batch applications and **analytics**.
- Common algorithms for best effort applications:
 - Round-robin.
 - First-Come-First-Serve (FCFS).
 - Shortest-Job-First (SJF).
 - Priority algorithms.

Cloud scheduling algorithms (2/2)

- Multimedia applications (e.g., audio and video streaming)
 - Have soft real-time constraints.
 - Require statistically guaranteed maximum delay and throughput.
- Real-time applications have hard real-time constraints.
- Scheduling algorithms for real-time applications:
 - Earliest Deadline First (**EDF**).
 - Rate Monotonic Algorithms (**RMA**).
- Algorithms for integrated scheduling of several classes of applications (best-effort, multimedia, real-time):
 - Resource Allocation/Dispatching (**RAD**) .
 - Rate-Based Earliest Deadline (**RBED**).

Scheduling Policies

- **First in, First out (FIFO)** →→ The tasks are scheduled for execution in the order of their arrival.
- **Earliest deadline first (EDF)** →→ The task with the earliest deadline is scheduled first.
- **Maximum workload derivative first (MWF)** →→ The tasks are scheduled in the order of their derivatives, the one with the highest derivative first. The number n of nodes assigned to the application is kept to a minimum.

Workload Partitioning Rules

- **Optimal Partitioning Rule (OPR)** → the workload is partitioned to ensure the earliest possible completion time and all tasks are required to complete at the same time.
 - The head node distributes sequentially the data to individual worker nodes.
 - Worker nodes start processing the data as soon as the transfer is complete.

- **Equal Partitioning Rule (EPR)** → assigns an equal workload to individual worker nodes.
 - The head node distributes sequentially the data to individual worker nodes.
 - Worker nodes start processing the data as soon as the transfer is complete.
 - The workload is partitioned in equal segments.

Scheduling MapReduce Applications

subject to deadlines

Four commonly referenced scheduling are:

- The default FIFO schedule
- The Fair Scheduler
- The Capacity Scheduler
- The Dynamic Proportional Scheduler

They are all based upon the assumptions:

1. The system is homogeneous
2. Load equipartition